

Towards Combining Statistical Relational Learning and Graph Neural Networks for Reasoning

Jian Tang

Mila-Quebec AI Institute

HEC Montreal

CIFAR AI Research Chair



Relational Data/Graphs are Ubiquitous

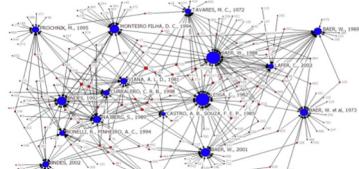
- Graphs: a general and flexible data structure to encode the relations between objects



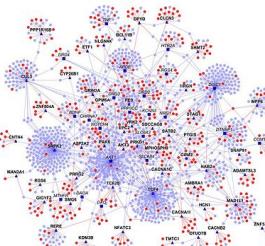
Social Graph



Road Graph

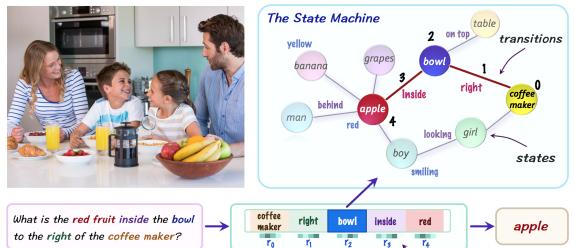
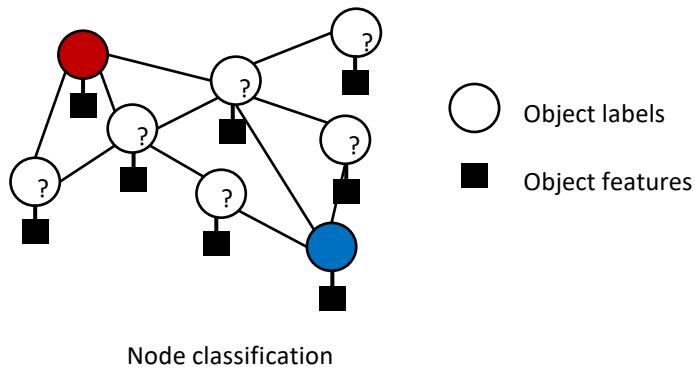


Citation Graph

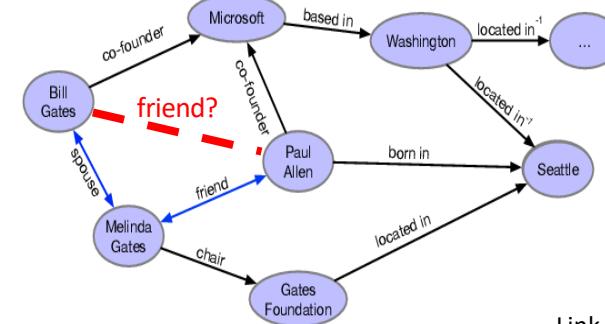


Protein-protein
Interaction Graph

Relational Prediction and Reasoning

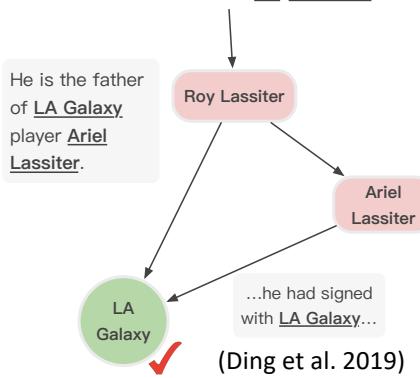


Visual relational reasoning
(Hudson et al. 2019)



Link prediction on knowledge graphs

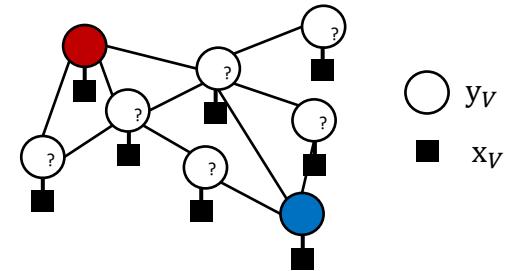
Q: What Cason, CA soccer team features the son of Roy Lassiter?



Multi-hop Question answering

Statistical Relational Learning

- Probabilistic graphical models for relational data
 - Markov Networks (Ross et al. 1980)
 - Conditional Random Fields (Lafferty et al. 2001)
 - Markov Logic Networks (Richardson and Domingos, 2006)
- Pros:
 - Captures uncertainty and domain knowledge
 - Collective inference
- Cons:
 - Limited model capacity
 - Inference is difficult

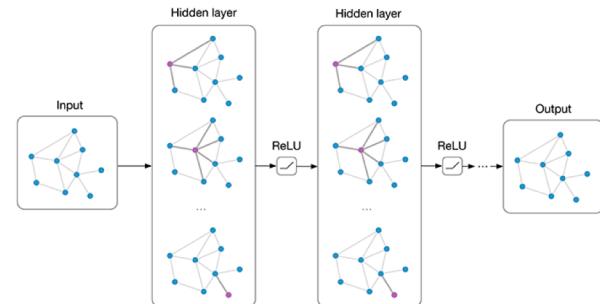
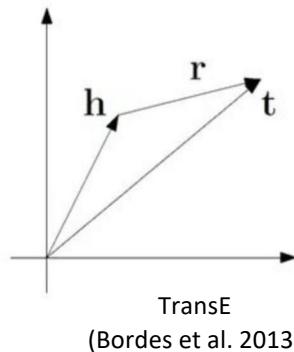
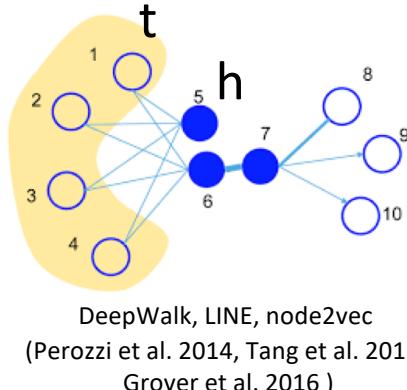


$$p(\mathbf{y}_V | \mathbf{x}_V) = \frac{1}{Z(\mathbf{x}_V)} \prod_{(i,j) \in E} \psi_{i,j}(\mathbf{y}_i, \mathbf{y}_j, \mathbf{x}_V)$$

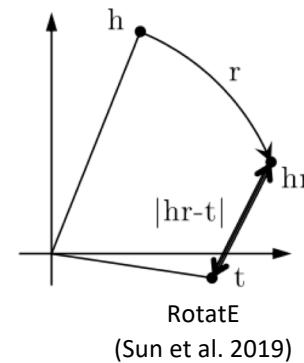
Figure: Conditional Random Fields

Graph Representation Learning

- Graph Neural Networks
 - Graph convolutional Networks (Kipf et al. 2016)
 - Graph attention networks (Veličković et al. 2017)
 - Neural message passing (Gilmer et al. 2017)
- Node Embedding and Knowledge Graph Embedding

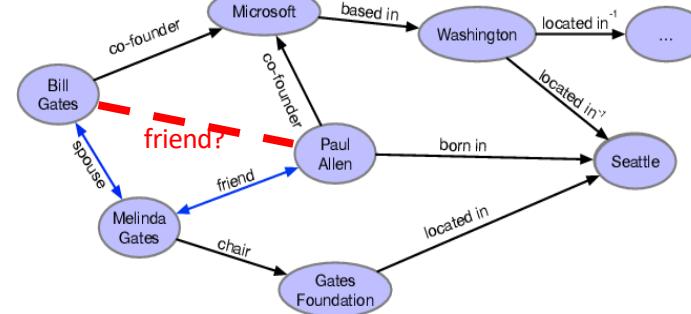


Graph convolutional Networks
(Kipf et al. 2016)



Link Prediction on Knowledge Graphs

- A set of facts $KG = \{(h, r, t)\}$ represented as triplets
 - (Bill_Gates, Co_Founder, Microsoft)
- A variety of applications
 - Question answering
 - Search
 - Recommender Systems
 - Natural language understanding
 - ...
- A fundamental problem: **predicting the missing facts by reasoning with existing facts**



Traditional Symbolic Logic-Rule based approaches

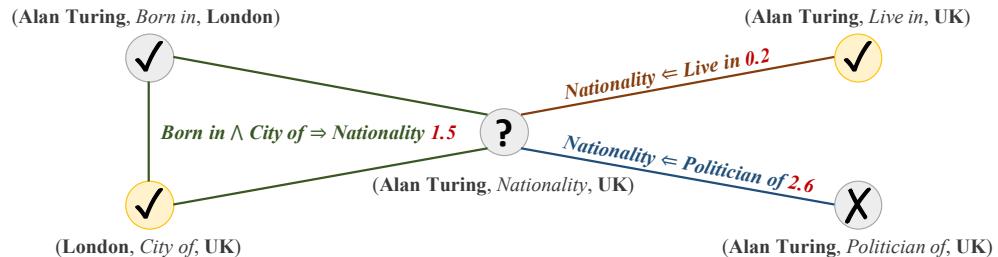
- Expert systems: hard logic rules
 - E.g., $\forall X, Y, \text{Husband}(X, Y) \Rightarrow \text{Wife}(Y, X)$
 - $\forall X, Y, \text{Live}(X, Y) \Rightarrow \text{Nationality}(X, Y)$
- Problematic as logic rules can be imperfect or contradictory
- We must handle the uncertainty of logic rules

Markov Logic Networks

(Richardson and Domingos, 2006)

- Combines first-order logic and probabilistic graphical models

- 0.2 $\text{Live}(X, Y) \Rightarrow \text{Nationality}(X, Y)$
- 2.6 $\text{Politician_of}(X, Y) \Rightarrow \text{Nationality}(X, Y)$
- 1.5 $\text{Born}(X, Y) \wedge \text{City_of}(Y, Z) \Rightarrow \text{Nationality}(X, Z)$



$$p(\mathbf{v}_O, \mathbf{v}_H) = \frac{1}{Z} \exp \left(\sum_{l \in L} w_l \sum_{g \in G_l} \mathbb{1}\{g \text{ is true}\} \right) = \frac{1}{Z} \exp \left(\sum_{l \in L} w_l n_l(\mathbf{v}_O, \mathbf{v}_H) \right)$$

\mathbf{v}_O : observed facts

\mathbf{v}_H : unobserved/hidden facts

w_l : weight of logic rule l

$n_l(\mathbf{v}_O, \mathbf{v}_H)$: number of true grounds of the logic rule l

Pros and Cons of Markov Logic Networks

- Pros
 - Effectively leverage domain knowledge with logic rules
 - Handle the uncertainty
- Limitation
 - Inference is difficult due to complicated graph structures
 - Recall is low since many facts are not covered by any logic rules

Knowledge Graph Embeddings

- Learning the entity and relation embeddings for predicting the missing facts (e.g., TransE, ComplEx, DisMult, RotatE)
- Defining the joint distribution of all the facts

$$p(\mathbf{v}_O, \mathbf{v}_H) = \prod_{(h,r,t) \in O \cup H} \text{Ber}(\mathbf{v}_{(h,r,t)} | f(\mathbf{x}_h, \mathbf{x}_r, \mathbf{x}_t)),$$

An example:

$$\text{Ber}(\mathbf{v}_{(h,r,t)} | f(\mathbf{x}_h, \mathbf{x}_r, \mathbf{x}_t)) = \sigma(\gamma - ||\mathbf{x}_h + \mathbf{x}_r - \mathbf{x}_t||) \quad \sigma \text{ is the sigmoid function, } \gamma \text{ is a fixed margin}$$

- Trained by treating V_O as positive facts and V_H as negative facts

Pros and Cons

- Pros
 - Can be effectively and efficiently trained by SGD
 - High recall of missing link prediction with entity and relation embeddings
- Cons
 - Hard to leverage domain knowledge (logic rules)

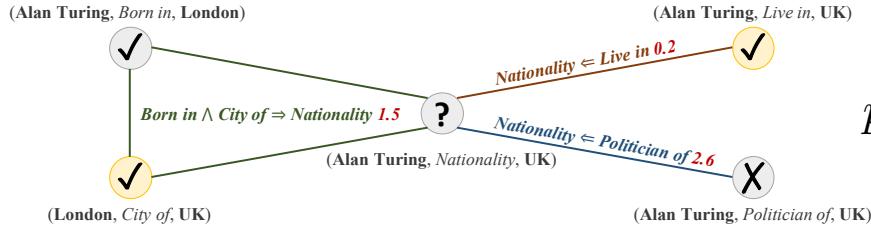
Probabilistic Logic Neural Networks for Reasoning (Qu and Tang, NeurIPS'19.)

- Towards combining Markov Logic Networks and knowledge graph embedding
 - Leverage logic rules and handling their uncertainty
 - Effective and efficient inference
- Define the joint distribution of facts with Markov Logic Network
- Optimization with variational EM
 - Parametrize the variational distribution with knowledge graph embedding methods

Meng Qu and Jian Tang. "Probabilistic Logic Neural Networks for Reasoning." To appear in NeurIPS'2019.

pLogicNet

- Define the joint distribution of facts with an MLN



$$p_w(\mathbf{v}_O, \mathbf{v}_H) = \frac{1}{Z} \exp \left(\sum_l w_l n_l(\mathbf{v}_O, \mathbf{v}_H) \right)$$

- Learning by maximizing the variational lower-bound of the log-likelihood of observed facts

$$\log p_w(\mathbf{v}_O) \geq \mathcal{L}(q_\theta, p_w) = \mathbb{E}_{q_\theta(\mathbf{v}_H)} [\log p_w(\mathbf{v}_O, \mathbf{v}_H) - \log q_\theta(\mathbf{v}_H)]$$

Inference

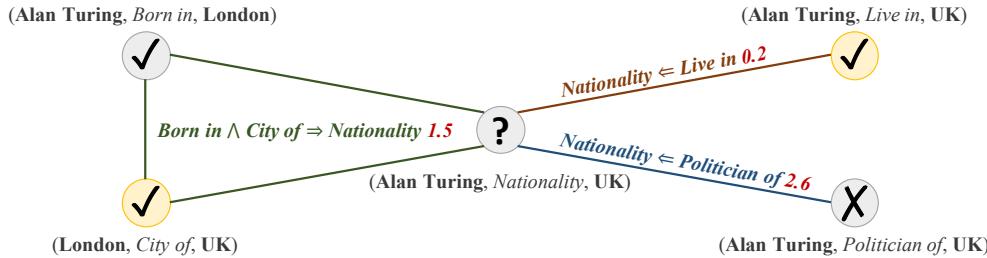
- Amortized mean-field variational inference
 - Use knowledge graph embedding model to parameterize the variational distribution

$$q_{\theta}(\mathbf{v}_H) = \prod_{(h,r,t) \in H} q_{\theta}(\mathbf{v}_{(h,r,t)}) = \prod_{(h,r,t) \in H} \text{Ber}(\mathbf{v}_{(h,r,t)} | f(\mathbf{x}_h, \mathbf{x}_r, \mathbf{x}_t)),$$

Learning

- Optimize pseudo-likelihood function
 - Update the weights of logic rules

$$\ell_{PL}(w) \triangleq \mathbb{E}_{q_\theta(\mathbf{v}_H)} \left[\sum_{h,r,t} \log p_w(\mathbf{v}_{(h,r,t)} | \mathbf{v}_{O \cup H \setminus (h,r,t)}) \right] = \mathbb{E}_{q_\theta(\mathbf{v}_H)} \left[\sum_{h,r,t} \log p_w(\mathbf{v}_{(h,r,t)} | \mathbf{v}_{\text{MB}(h,r,t)}) \right],$$



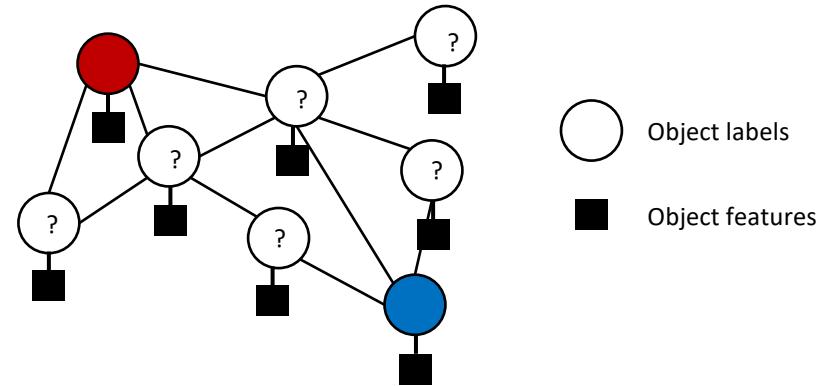
Performance of Link Prediction

- **Datasets:** benchmark knowledge graphs
 - FB15K, WN18, FB15K-237, WN18-RR
- Logic rules:
 - Composition rules (e.g., Father of Father is GrandFather)
 - Inverse rules (e.g., Husband and Wife)
 - Symmetric rules (e.g., Similar)
 - Subrelation rules (e.g., Man => Person)

Category	Algorithm	FB15k					WN18				
		MR	MRR	H@1	H@3	H@10	MR	MRR	H@1	H@3	H@10
KGE	TransE [3]	40	0.730	64.5	79.3	86.4	272	0.772	70.1	80.8	92.0
	DistMult [17]	42	0.798	-	-	89.3	655	0.797	-	-	94.6
	Hole [26]	-	0.524	40.2	61.3	73.9	-	0.938	93.0	94.5	94.9
	ComplEx [41]	-	0.692	59.9	75.9	84.0	-	0.941	93.6	94.5	94.7
	ConvE [8]	51	0.657	55.8	72.3	83.1	374	0.943	93.5	94.6	95.6
Rule-based	BLP [7]	415	0.242	15.1	26.9	42.4	736	0.643	53.7	71.7	83.0
	MLN [32]	352	0.321	21.0	37.0	55.0	717	0.657	55.4	73.1	83.9
Hybrid	RUGE [15]	-	0.768	70.3	81.5	86.5	-	-	-	-	-
	NNE-AER [9]	-	0.803	76.1	83.1	87.4	-	0.943	94.0	94.5	94.8
Ours	pLogicNet	33	0.792	71.4	85.7	90.1	255	0.832	71.6	94.4	95.7
	pLogicNet*	33	0.844	81.2	86.2	90.2	254	0.945	93.9	94.7	95.8

Semi-supervised Object Classification

- Given $G = (V, E, \mathbf{x}_V)$
 - $V = V_L \cup V_U$: objects/nodes
 - E : edges
 - \mathbf{x}_V : object features
- Give some labeled objects V_L , we want to infer the labels of the rest of objects V_U



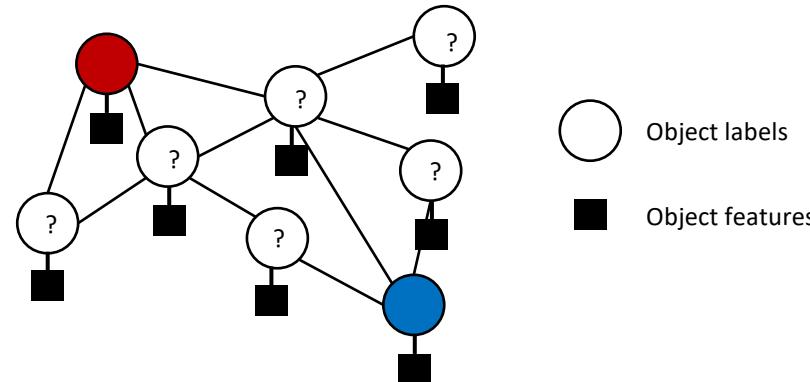
GMNN: Graph Markov Neural Networks (Qu, Bengio, and Tang, ICML'19)

- Combining conditional random fields and graph neural networks
 - Learning effective node representations
 - Modeling the label dependencies of nodes
- Model the joint distribution of object labels \mathbf{y}_V conditioned on object attributes \mathbf{x}_V , i.e., $p_\phi(\mathbf{y}_V | \mathbf{x}_V)$ with CRFs
 - Optimization with Pseudolikelihood Variational-EM

$$\log p_\phi(\mathbf{y}_L | \mathbf{x}_V) \geq \mathbb{E}_{q_\theta(\mathbf{y}_U | \mathbf{x}_V)} [\log p_\phi(\mathbf{y}_L, \mathbf{y}_U | \mathbf{x}_V) - \log q_\theta(\mathbf{y}_U | \mathbf{x}_V)]$$

Overall Optimization Procedure

- Two Graph Neural Networks Collaborate with each other
 - p_ϕ : learning network, modeling the label dependency
 - q_θ : inference network, learning the object representations
- q_θ infer the labels of unlabeled objects trained with supervision from p_ϕ and labeled objects
- p_ϕ is trained with a fully labeled graph, where the unlabeled objects are labeled by q_θ



Take Away

- Relational reasoning is important to a variety of applications
 - Node classification, link prediction on knowledge graphs, question answering
- Towards combining two learning frameworks
 - Statistical Relational Learning
 - Graph Representation Learning
- Looking forward
 - Combining deep learning and symbolic reasoning systems
 - Incorporating common sense knowledge, handling uncertainty, and maybe automatically learn the logic rules.

Questions?
Email: jian.tang@hec.ca

Results on FB15k-237 and WN18RR

Category	Algorithm	FB15k-237					WN18RR				
		MR	MRR	H@1	H@3	H@10	MR	MRR	H@1	H@3	H@10
KGE	TransE [3]	181	0.326	22.9	36.3	52.1	3410	0.223	1.3	40.1	53.1
	DistMult [17]	254	0.241	15.5	26.3	41.9	5110	0.43	39	44	49
	ComplEx [41]	339	0.247	15.8	27.5	42.8	5261	0.44	41	46	51
	ConvE [8]	244	0.325	23.7	35.6	50.1	4187	0.43	40	44	52
Rule-based	BLP [7]	1985	0.092	6.2	9.8	15.0	12051	0.254	18.7	31.3	35.8
	MLN [32]	1980	0.098	6.7	10.3	16.0	11549	0.259	19.1	32.2	36.1
Ours	pLogicNet	173	0.330	23.1	36.9	52.8	3436	0.230	1.5	41.1	53.1
	pLogicNet*	173	0.332	23.7	36.7	52.4	3408	0.441	39.8	44.6	53.7

GMNN: Graph Markov Neural Networks

- Model the joint distribution of object labels \mathbf{y}_V conditioned on object attributes \mathbf{x}_V , i.e., $p_\phi(\mathbf{y}_V | \mathbf{x}_V)$
- Learning the model parameters ϕ by maximizing the lower-bound of log-likelihood of the observed data, $\log p_\phi(\mathbf{y}_L | \mathbf{x}_V)$

$$\begin{aligned} \log p_\phi(\mathbf{y}_L | \mathbf{x}_V) &\geq \\ \mathbb{E}_{q_\theta(\mathbf{y}_U | \mathbf{x}_V)} [\log p_\phi(\mathbf{y}_L, \mathbf{y}_U | \mathbf{x}_V) - \log q_\theta(\mathbf{y}_U | \mathbf{x}_V)] \end{aligned}$$

Optimization with Pseudolikelihood Variational-EM

- E-step: fix p_ϕ and update the variational distribution $q_\theta(\mathbf{y}_U | \mathbf{x}_V)$ to approximate the true posterior distribution $p_\phi(\mathbf{y}_U | \mathbf{y}_L, \mathbf{x}_V)$.
- M-step: fix q_θ and update p_ϕ to maximize the lower bound

$$\ell(\phi) = \mathbb{E}_{q_\theta(\mathbf{y}_U | \mathbf{x}_V)} [\log p_\phi(\mathbf{y}_L, \mathbf{y}_U | \mathbf{x}_V)]$$

- Directly optimize the joint likelihood is difficult due to the partition function in p_ϕ , instead we optimize the pseudolikelihood function

$$\ell_{PL}(\phi) \triangleq \mathbb{E}_{q_\theta(\mathbf{y}_U | \mathbf{x}_V)} \left[\sum_{n \in V} \log p_\phi(\mathbf{y}_n | \mathbf{y}_{V \setminus n}, \mathbf{x}_V) \right]$$

$$= \mathbb{E}_{q_\theta(\mathbf{y}_U | \mathbf{x}_V)} \left[\sum_{n \in V} \log p_\phi(\mathbf{y}_n | \mathbf{y}_{\text{NB}(n)}, \mathbf{x}_V) \right]$$

Inference/E-step: approximate $p_\phi(\mathbf{y}_U | \mathbf{y}_L, \mathbf{x}_V)$

- Approximate it with variational distribution $q_\theta(\mathbf{y}_U | \mathbf{x}_V)$. Specifically we use mean-field method:

$$q_\theta(\mathbf{y}_U | \mathbf{x}_V) = \prod_{n \in U} q_\theta(\mathbf{y}_n | \mathbf{x}_V).$$

- We parametrize each variational distribution with a Graph Neural Network

$$q_\theta(\mathbf{y}_n | \mathbf{x}_V) = \text{Cat}(\mathbf{y}_n | \text{softmax}(W_\theta \mathbf{h}_{\theta, n}))$$

Object representations learned by GNN

Learning/M-step:

- The log-pseudo likelihood:

$$\begin{aligned}\ell_{PL}(\phi) &\triangleq \mathbb{E}_{q_\theta(\mathbf{y}_U | \mathbf{x}_V)} \left[\sum_{n \in V} \log p_\phi(\mathbf{y}_n | \mathbf{y}_{V \setminus n}, \mathbf{x}_V) \right] \\ &= \mathbb{E}_{q_\theta(\mathbf{y}_U | \mathbf{x}_V)} \left[\sum_{n \in V} \log p_\phi(\mathbf{y}_n | \mathbf{y}_{\text{NB}(n)}, \mathbf{x}_V) \right]\end{aligned}$$

- According to the inference, only the $p_\phi(\mathbf{y}_n | \mathbf{y}_{\text{NB}(n)}, \mathbf{x}_V)$ is required
- Parametrize $p_\phi(\mathbf{y}_n | \mathbf{y}_{\text{NB}(n)}, \mathbf{x}_V)$ with another GCN

$$p_\phi(\mathbf{y}_n | \mathbf{y}_{\text{NB}(n)}, \mathbf{x}_V) = \text{Cat}(\mathbf{y}_n | \text{softmax}(W_\phi \mathbf{h}_{\phi, n}))$$

Overall Optimization Procedure

- Two Graph Neural Networks Collaborate with each other
 - p_ϕ : learning network, modeling the label dependency
 - q_θ : inference network, learning the object representations
- q_θ infer the labels of unlabeled objects trained with supervision from p_ϕ and labeled objects
- p_ϕ is trained with a fully labeled graph, where the unlabeled objects are labeled by q_θ

